

SK LEARN:- package should be installed.

#unwanted data should be removed first. #positive correlation "(+1)",BLUE COLOUR. #negative correlation "(-1)",BLACK COLOUR.

3STEPS:-

1)READ THE DATA. => data=pd.read_csv("_____.csv") 2)DESCRIBE THE DATA. => data.describe() 3)INFO. => data.info()

RMSE:-Root Mean Square Error.

```
In [151]: import pandas as pd
```

```
In [152]: data=pd.read_csv("/home/placement/Desktop/ramaraju/fiat500.csv")
```

In [153]: data

Out[153]:

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

In [154]: data.describe()

Out[154]:

	ID	engine_power	age_in_days	km	previous_owners	lat	lon	price
count	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
mean	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.563428	8576.003901
std	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.328190	1939.958641
min	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.245400	2500.000000
25%	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.505090	7122.500000
50%	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.869260	9000.000000
75%	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.769040	10000.000000
max	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.365520	11100.000000

In [155]: data.info

```
Out[155]: <bound method DataFrame.info of
0      1 lounge      51      ID      model  engine_power  age_in_days      km  previous_owners  \
1      2   pop      51     882    25000                1
2      3  sport      74    1186    32500                1
3      4 lounge      51   4658   142228                1
4      5   pop      73   2739   160000                1
...    ...    ...    ...    ...    ...    ...
1533  1534  sport      51   3074   106880                1
1534  1535 lounge      74   3712   115280                1
1535  1536   pop      51   3835   112000                1
1536  1537 lounge      51   2223    60457                1
1537  1538   pop      51   2557    80750                1
1537  1538   pop      51   1766    54276                1

      lat      lon  price
0  44.907242  8.611560  8900
1  45.666359  12.241890  8800
2  45.503300  11.417840  4200
3  40.633171  17.634609  6000
4  41.903221  12.495650  5700
...    ...    ...    ...
1533  45.069679  7.704920  5200
1534  45.845692  8.666870  4600
1535  45.481541  9.413480  7500
1536  45.000702  7.682270  5990
1537  40.323410  17.568270  7900

[1538 rows x 9 columns]>
```

```
In [156]: data1=data.drop(["lat","lon","ID"],axis=1)
data1
```

```
Out[156]:
```

	model	engine_power	age_in_days	km	previous_owners	price
0	lounge	51	882	25000	1	8900
1	pop	51	1186	32500	1	8800
2	sport	74	4658	142228	1	4200
3	lounge	51	2739	160000	1	6000
4	pop	73	3074	106880	1	5700
...
1533	sport	51	3712	115280	1	5200
1534	lounge	74	3835	112000	1	4600
1535	pop	51	2223	60457	1	7500
1536	lounge	51	2557	80750	1	5990
1537	pop	51	1766	54276	1	7900

1538 rows × 6 columns

```
In [157]: data1=pd.get_dummies(data1)
```

```
In [158]: data1
```

```
Out[158]:
```

	engine_power	age_in_days	km	previous_owners	price	model_lounge	model_pop	model_sport
0	51	882	25000	1	8900	1	0	0
1	51	1186	32500	1	8800	0	1	0
2	74	4658	142228	1	4200	0	0	1
3	51	2739	160000	1	6000	1	0	0
4	73	3074	106880	1	5700	0	1	0
...
1533	51	3712	115280	1	5200	0	0	1
1534	74	3835	112000	1	4600	1	0	0
1535	51	2223	60457	1	7500	0	1	0
1536	51	2557	80750	1	5990	1	0	0
1537	51	1766	54276	1	7900	0	1	0

1538 rows × 8 columns

```
In [159]: y=data1['price']  
x=data1.drop('price',axis=1)
```

In [160]:

y

Out[160]:

0	8900
1	8800
2	4200
3	6000
4	5700

...

1533	5200
1534	4600
1535	7500
1536	5990
1537	7900

Name: price, Length: 1538, dtype: int64

In [161]:

x

Out[161]:

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
0	51	882	25000	1	1	0	0
1	51	1186	32500	1	0	1	0
2	74	4658	142228	1	0	0	1
3	51	2739	160000	1	1	0	0
4	73	3074	106880	1	0	1	0
...
1533	51	3712	115280	1	0	0	1
1534	74	3835	112000	1	1	0	0
1535	51	2223	60457	1	0	1	0
1536	51	2557	80750	1	1	0	0
1537	51	1766	54276	1	0	1	0

1538 rows × 7 columns

```
In [162]: #!/pip install scikit-learn
#just now installed .so, i kept in command.no need to use again once installed.
```

```
In [163]: #calling a function to split.
#split enter data into ->67% training, ->23% testing.
from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33,random_state=42)
```

```
In [164]: x_train.head(5)
```

```
Out[164]:
```

	engine_power	age_in_days	km	previous_owners	model_lounge	model_pop	model_sport
527	51	425	13111	1	1	0	0
129	51	1127	21400	1	1	0	0
602	51	2039	57039	1	0	1	0
331	51	1155	40700	1	1	0	0
323	51	425	16783	1	1	0	0

```
In [165]: y_train.head(5)
```

```
Out[165]: 527    9990
129    9500
602    7590
331    8750
323    9100
Name: price, dtype: int64
```

```
In [166]: from sklearn.linear_model import LinearRegression
reg=LinearRegression()
reg.fit(x_train,y_train)
```

```
Out[166]: LinearRegression()
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [167]: ypred=reg.predict(x_test)
```

```
In [168]: ypred
```

```
Out[168]: array([ 5867.6503378 ,  7133.70142341,  9866.35776216,  9723.28874535,  
 10039.59101162,  9654.07582608,  9673.14563045, 10118.70728123,  
  9903.85952664,  9351.55828437, 10434.34963575,  7732.26255693,  
  7698.67240131,  6565.95240435,  9662.90103518, 10373.20344286,  
  9599.94844451,  7699.34400418,  4941.33017994, 10455.2719478 ,  
 10370.51555682, 10391.60424404,  7529.06622456,  9952.37340054,  
  7006.13845729,  9000.1780961 ,  4798.36770637,  6953.10376491,  
  7810.39767825,  9623.80497535,  7333.52158317,  5229.18705519,  
  5398.21541073,  5157.65652129,  8948.63632836,  5666.62365159,  
  9822.1231461 ,  8258.46551788,  6279.2040404 ,  8457.38443276,  
  9773.86444066,  6767.04074749,  9182.99904787, 10210.05195479,  
  8694.90545226, 10328.43369248,  9069.05761443,  8866.7826029 ,  
  7058.39787506,  9073.33877162,  9412.68162121, 10293.69451263,  
 10072.49011135,  6748.5794244 ,  9785.95841801,  9354.09969973,  
  9507.9444386 , 10443.01608254,  9795.31884316,  7197.84932877,  
 10108.31707235,  7009.6597206 ,  9853.90699412,  7146.87414965,  
  6417.69133992,  9996.97382441,  9781.18795953,  8515.83255277,  
  8456.30006203,  6499.76668237,  7768.57829985,  6832.86406122,  
  8347.96113362, 10439.02404036,  7356.43463051,  8562.56562053,  
  8820.78555100, 10025.02571520,  7270.77100022,  8411.45004000])
```

```
In [169]: from sklearn.metrics import r2_score  
r2_score(y_test,ypred)
```

```
Out[169]: 0.8415526986865394
```

```
In [170]: #RMSE  
from sklearn.metrics import mean_squared_error  
mean_squared_error(ypred,y_test)
```

```
Out[170]: 581887.727391353
```



```
In [174]: Results=pd.DataFrame(columns=['price','predicted'])
Results['price']=y_test
Results['predicted']=ypred
Results=Results.reset_index()
Results['ID']=Results.index
Results.head(15)
```

Out[174]:

	price	predicted	ID
481	7900	5867.650338	481
76	7900	7133.701423	76
1502	9400	9866.357762	1502
669	8500	9723.288745	669
1409	9700	10039.591012	1409
1414	9900	9654.075826	1414
1089	9900	9673.145630	1089
1507	9950	10118.707281	1507
970	10700	9903.859527	970
1198	8999	9351.558284	1198
1088	9890	10434.349636	1088
576	7990	7732.262557	576
965	7380	7698.672401	965
1488	6800	6565.952404	1488
1432	8900	9662.901035	1432

```
In [173]: #Results.plot()
```

```
In [175]: Results["predicted value"]=Results.apply(lambda row:row.price-row.predicted,axis=1)
Results
```

```
Out[175]:
```

	price	predicted	ID	predicted value
481	7900	5867.650338	481	2032.349662
76	7900	7133.701423	76	766.298577
1502	9400	9866.357762	1502	-466.357762
669	8500	9723.288745	669	-1223.288745
1409	9700	10039.591012	1409	-339.591012
...
291	10900	10032.665135	291	867.334865
596	5699	6281.536277	596	-582.536277
1489	9500	9986.327508	1489	-486.327508
1436	6990	8381.517020	1436	-1391.517020
575	10900	10371.142553	575	528.857447

508 rows × 4 columns

```
In [ ]:
```